

2019 Deep Learning and Practice

Lab 7 – Temporal Difference Learning

0756110 李東霖

May 29, 2019

1 Introduction

Play 2048 through Temporal Difference Learning, a kind of Reinforcement Learning.

Requirements as follows:

- Understand TD(0)
- Implement TD(0)
- Train 2048 agent with before-state and after-state

1.1 Game Environment – 2048

There are UP, DOWN, LEFT, RIGHT four action in 2048 game. The reward of game is the value of new tile when two tiles are combined. For example, 4 and 4 can be combined as 8 and then the reward is 8.

2 Experiment setup

2.1 Temporal Difference (0)

Temporal Difference Learning is model-free algorithm thus it need estimate future reward by episodes of experience. But if like Monte-Carlo learns from complete episodes, its learning is slow. To avoid that problem, TD learns from incomplete episodes, also means learning from next few states. We can see equation of TD(0) as follow.

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t)) \quad (1)$$

If $V(S_{t+1})$ is actual value, TD is unbiased estimate. But $V(S_{t+1})$ is also from estimated. Therefore TD(0) is biased estimate. This is different from MC algorithm.

2.2 Before state

The game 2048 has a random state transition after doing action.

$$s_t \xrightarrow{a_t} s'_t \xrightarrow{\text{random popup}} s''_t, s_{t+1} = s''_t \quad (2)$$

Before state means that estimate s_t with s''_t . So use equation as follows:

$$V(s_t) \leftarrow V(s_t) + \alpha(r + \gamma V(s''_t) - V(s_t)) \quad (3)$$

And find best action to use as follows:

$$\arg \max r + \sum_{s' \in P(s'|s,a), s'' \in P(s''|s')} P(s''|s') V(s'') \quad (4)$$

It means that find expectation of all possible s'' and choose maximum expectation.

2.3 After state

Contrary to before state, after state estimate s'_t with s'_{t+1} . So use equation as follows:

$$V(s'_t) \leftarrow V(s'_t) + \alpha(r + \gamma V(s'_{t+1}) - V(s'_t)) \quad (5)$$

And find best action to use as follows:

$$\arg \max r + V(s'_t) \quad (6)$$

It means ignore random popup, only to estimate state after action.

2.4 My implementation (python)

2.4.1 Game environment

Game environment is reference from <https://github.com/moporgic/2048-Demo-Python/>. I added `rotate`, `end` and `allpopup` functions. The `rotate` function can rotate whole game board. The `end` function can judge game is end or not. The `allpopup` function can output all possible game after game pop up with probability.

```
class board:
    """simple implementation of 2048 puzzle"""

    def __init__(self, tile = None, max_number=15):
        self.tile = tile if tile is not None else [0] * 16
        self.max_num = max_number

    def __str__(self):
        state = '+' + '-' * 24 + '+\n'
        for row in [self.tile[r:r + 4] for r in range(0, 16, 4)]:
```

```

        state += ('|' + ''.join('{0:6d}'.format((1 << t) & -2) for t in
        ↪ row) + '|\\n')
state += '+' + '-' * 24 + '+'
return state

def mirror(self):
    return board([self.tile[r + i] for r in range(0, 16, 4) for i in
    ↪ reversed(range(4))])

def transpose(self):
    return board([self.tile[r + i] for i in range(4) for r in range(0,
    ↪ 16, 4)])

def rotate(self):
    return board([self.tile[4*(3-(i%4)) + (i//4)] for i in range(16)])

def left(self):
    move, score = [], 0
    for row in [self.tile[r:r+4] for r in range(0, 16, 4)]:
        row, buf = [], [t for t in row if t]
        while buf:
            if len(buf) >= 2 and buf[0] is buf[1]:
                buf = buf[1:]
                buf[0] += 1
                score += 1 << buf[0]
            row += [buf[0]]
            buf = buf[1:]
        move += row + [0] * (4 - len(row))
    return board(move), score if move != self.tile else -1

def right(self):
    move, score = self.mirror().left()
    return move.mirror(), score

def up(self):
    move, score = self.transpose().left()
    return move.transpose(), score

def down(self):
    move, score = self.transpose().right()
    return move.transpose(), score

def popup(self):
    tile = self.tile[:]
    empty = [i for i, t in enumerate(tile) if not t]

```

```

tile[random.choice(empty)] = random.choice([1] * 9 + [2])
return board(tile)

def allpopup(self):
    tile = self.tile[:]
    empty = [i for i, t in enumerate(tile) if not t]
    boards = []
    for i in empty:
        for n in [1] * 9 + [2]:
            tmp = tile.copy()
            tmp[i] = n
            boards.append(board(tmp))
    return boards

def end(self):
    tile = self.tile[:]
    empty = [i for i, t in enumerate(tile) if not t]

    count_max_num = np.count_nonzero(self.max_num == np.array(tile))
    return len(empty) == 0 or count_max_num > 0

```

2.4.2 Tuple network

I implemented a class to perform N-Tuple Network with all possible isomorphic pattern.

```

def find_isomorphic_pattern(pattern):
    a = board(list(range(16)))

    isomorphic_pattern = []
    for i in range(8):
        if (i >= 4):
            b = board( a.mirror().tile )
        else:
            b = board( a.tile )
        for _ in range(i%4):
            b = b.rotate()
        isomorphic_pattern.append(np.array(b.tile)[pattern])

    return isomorphic_pattern

class TuplesNet():
    def __init__(self, pattern, maxnum):
        self.V = np.zeros(( [maxnum]*len(pattern)))
        self.pattern = pattern

```

```

        self.isomorphic_pattern = find_isomorphic_pattern(self.pattern)

    def getState(self, tile):
        return [tuple(np.array(tile)[p]) for p in self.isomorphic_pattern]

    def getValue(self, tile):
        S = self.getState(tile)

        V = [self.V[s] for s in S]

        # sum all value from isomorphic pattern
        V = sum(V)

        return V

    def setValue(self, tile, v, reset=False):
        S = self.getState(tile)

        v /= len(self.isomorphic_pattern)
        V = 0.0
        for s in S:
            if not reset:
                # update value to isomorphic pattern
                self.V[s] += v
            else:
                # reset value to isomorphic pattern
                self.V[s] = v

        V += self.V[s]
        return V

```

2.4.3 Agent

First, create agent with specific patterns. The agent use patterns to create `TuplesNet`.

```

class Agent():
    def __init__(self, patterns, maxnum):
        self.Tuples = []
        for p in patterns:
            self.Tuples.append(TuplesNet(p, maxnum))
        self.metrics = []
        # if True, use after-state. Otherwise use before-state
        self.after = True

```

Second, integrate agent with multiple TuplesNet through `getValue` and `setValue`. That two function let agent store value based on state.

```
def getValue(self, tile):
    return sum([t.getValue(tile) for t in self.Tuples])

def setValue(self, tile, v, reset=False):
    v /= len(self.Tuples)
    V = 0.0
    for t in self.Tuples:
        V += t.setValue(tile, v, reset)
    return V
```

Third, implement `evaluate` and `learn`. I use `self.after` to decide which method (before or after) to use. That two function let agent make decision through value function.

```
# get all s' and reward in next_games
def evaluate(self, next_games):
    # TD(0)-after
    if self.after:
        #  $r + V(s')$ 
        return [ng[1] + self.getValue(ng[0].tile) for ng in next_games]
    # TD(0)-before
    else:
        #  $r + \sum P(s''/s')V(s'')$ 
        rs = []
        for ng in next_games:
            all_v = [ self.getValue(nng.tile) for nng in
                ng[0].allpopup() ]
            if len(all_v) == 0:
                v = 0
            else:
                v = sum(all_v) / len(all_v)
            rs.append(ng[1] + v)
        return rs

def learn(self, records, lr):
    # learn from end to begin
    # records = [end .... begin]
    # (s, a, r, s', s'')

    # TD(0)-after
    if self.after:
```

```

exact = 0.0
for s, a, r, s_, s__ in records:
    #  $V(s') = V(s') + \alpha (r_{next} + V(s'_{next}) - V(s'))$ 
    error = exact - self.getValue(s_)
    exact = r + self.setValue(s_, lr*error)
# TD(0)-before
else:
    #exact = self.getValue(records[0][4])
    exact = 0.0
    for s, a, r, s_, s__ in records:
        #  $V(s) = V(s) + \alpha (r + V(s'') - V(s))$ 
        error = r + exact - self.getValue(s)
        exact = r + self.setValue(s, lr*error)

```

Finally, I need a train procedure to run TD(0) algorithm.

```

def train(self, epoch_size, lr=0.1, showsize=1000):
    start_epoch = len(self.metrics)
    for epoch in range(start_epoch, epoch_size):
        # init score and env (2048)
        score = 0.0
        game = board().popup().popup()
        records = []
        while True:
            # choose action
            next_games = [game.up(), game.down(), game.left(),
                          ↪ game.right()]
            action = np.argmax(self.evaluate(next_games))

            # do action
            # s'
            next_game, reward = next_games[action]

            # save record (s, a, r, s')
            # records.insert(0, (game.tile, action, reward,
            ↪ next_game.tile) )

            # if game is same as before, end game
            if game.end():
                break

            # s''
            next_game_after = next_game.popup()

```

```

score += reward

# save record (s, a, r, s', s'')
records.insert(0, (game.tile, action, reward,
    ↪ next_game.tile, next_game_after.tile) )

# s = s'' update state
game = next_game_after

#self.learn(records, lr / len(self.Tuples))
self.learn(records, lr)

# store score, game len, end game board
self.metrics.append( (score, len(records), game.tile) )
if (epoch+1) % showsize == 0:
    clear_output(wait=True)
    self.showStatistic(epoch+1, showsize)

```

And also implement some utility function to make convenient.

```

def showStatistic(self, epoch, unit, show=True):
    metrics = np.array(self.metrics[epoch-unit:epoch])

    # get average score
    score_mean = np.mean(metrics[:, 0])
    # get max score
    score_max = np.max(metrics[:, 0])

    if show:
        print('{:<8d}mean = {:<8.0f} max = {:<8.0f}'.format(epoch,
            ↪ score_mean, score_max))

    if (metrics.shape[1] < 3):
        return score_mean, score_max

    # all end game board
    end_games = metrics[:, 2]

    reach_nums = np.array([1<<max(end) & -2 for end in end_games])

    if show:
        print('\n')

    score_stat = []

```

```

for num in np.sort(np.unique(reach_nums)):
    # count how many game over this num
    reaches = np.count_nonzero(reach_nums >= num)
    reaches = (reaches*100)/len(metrics)

    # count how many game end at this num
    ends = np.count_nonzero(reach_nums == num)
    ends = (ends*100)/len(metrics)

    if show:
        print('{:<5d}  {:3.1f} % ({:3.1f} %)' .format(num, reaches,
            ↪ ends) )

    score_stat.append( (num, reaches, ends) )

score_stat = np.array(score_stat)

return score_mean, score_max, score_stat

# use current state of game, return next game and action
def play(self, game):
    next_games = [game.up(), game.down(), game.left(), game.right()]
    action = np.argmax(self.evaluate(next_games))

    next_game, reward = next_games[action]
    return next_game, reward, ['up', 'down', 'left', 'right'][action]

```

2.4.4 Main

```

MAX_NUM = 15 # 1<<15 == 32768
PATTERNS = [
    [0,1,2,3,4,5],
    [4,5,6,7,8,9],
    [0,1,2,4,5,6],
    [4,5,6,8,9,10]
]
random.seed(756110)
agent = Agent(PATTERNS, MAX_NUM)
agent.train(100000)

showCurve(agent.metrics, size=1000)
showWinRate(agent.metrics, [1024, 2048, 4096])

```

3 Experimental results

3.1 Score curves in Training




Figure 1: max/min/mean score curves every 1000 epoch

3.2 Win rate curves in Training

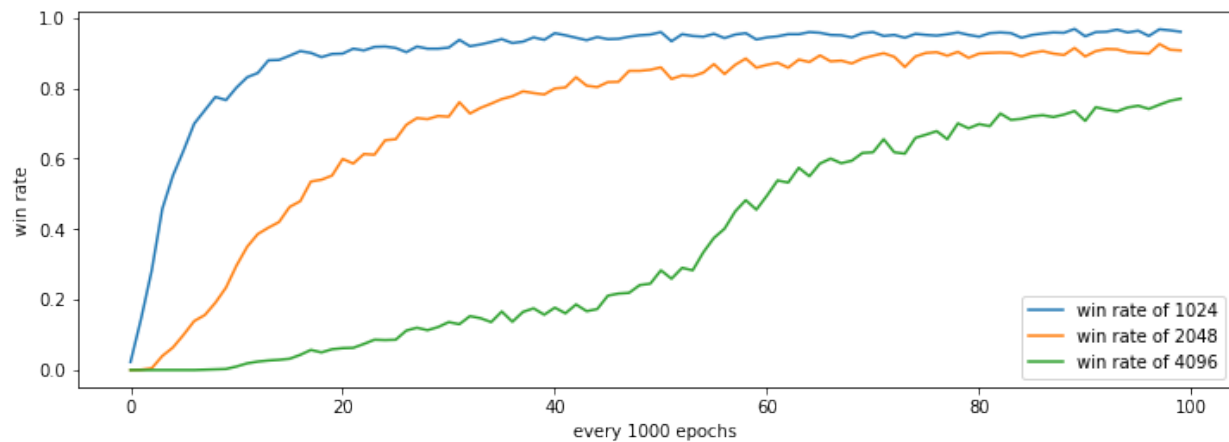


Figure 2: 1024 2048 4096 win rate curve every 1000 epoch

3.3 Agent performance

I use this agent to play 2048 with 1000 times.

NUM	WIN RATE	END RATE
64	100.0 %	0.1 %
128	99.9 %	0.1 %
256	99.8 %	1.7 %
512	98.1 %	2.0 %
1024	96.1 %	4.7 %
2048	91.4 %	11.9 %
4096	79.5 %	77.7 %
8192	1.8 %	1.8 %

Table 1: agent performance

The **NUM** columns means which number on 2048 tile. The **WIN RATE** columns means how many game can reach number on tile. The **END RATE** columns means how many max of number on tile at end game.

4 Discussion

4.1 Before state v.s. After state

My before state is very bad than after state. I think before state need all possible s'' . But all s'' usually are wrong value in the begin. And number of state is huge. The before state has slow learning because of that reason.

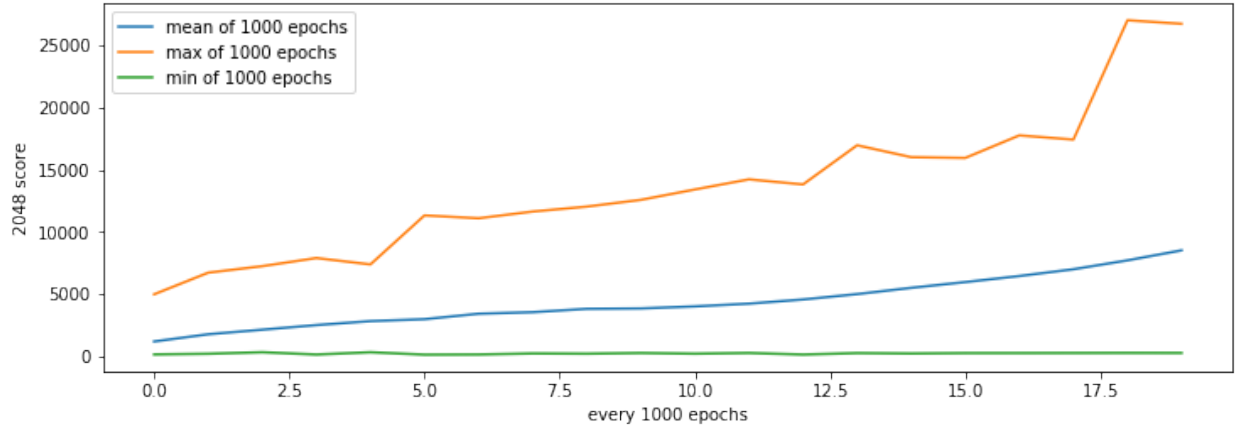


Figure 3: [Before state] max/min/mean score curves every 1000 epoch