

2019 Deep Learning and Practice  
Lab 8 – Deep Q-Learning and Deep Deterministic  
Policy Gradient

0756110 李東霖

June 5, 2019

## 1 DQN's episode rewards in training

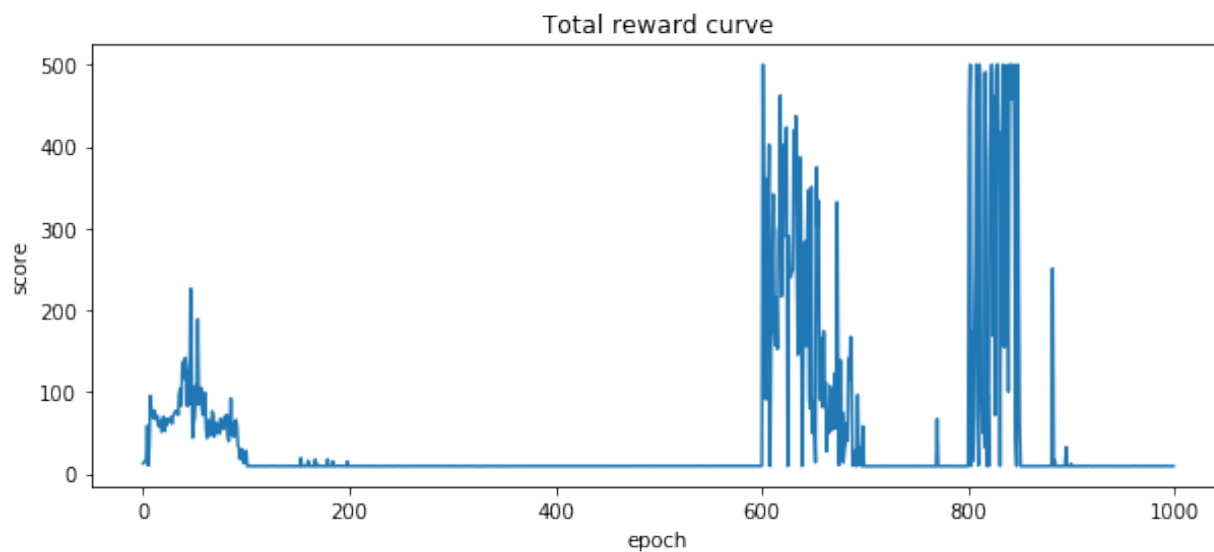


Figure 1: DQN average of total reward in 100 testing

## 2 DDPG's episode rewards in training




Figure 2: DDPG average of total reward in 100 testing

## 3 DQN structure and loss

I use two fully connected layer to build Deep Q Network as follows:

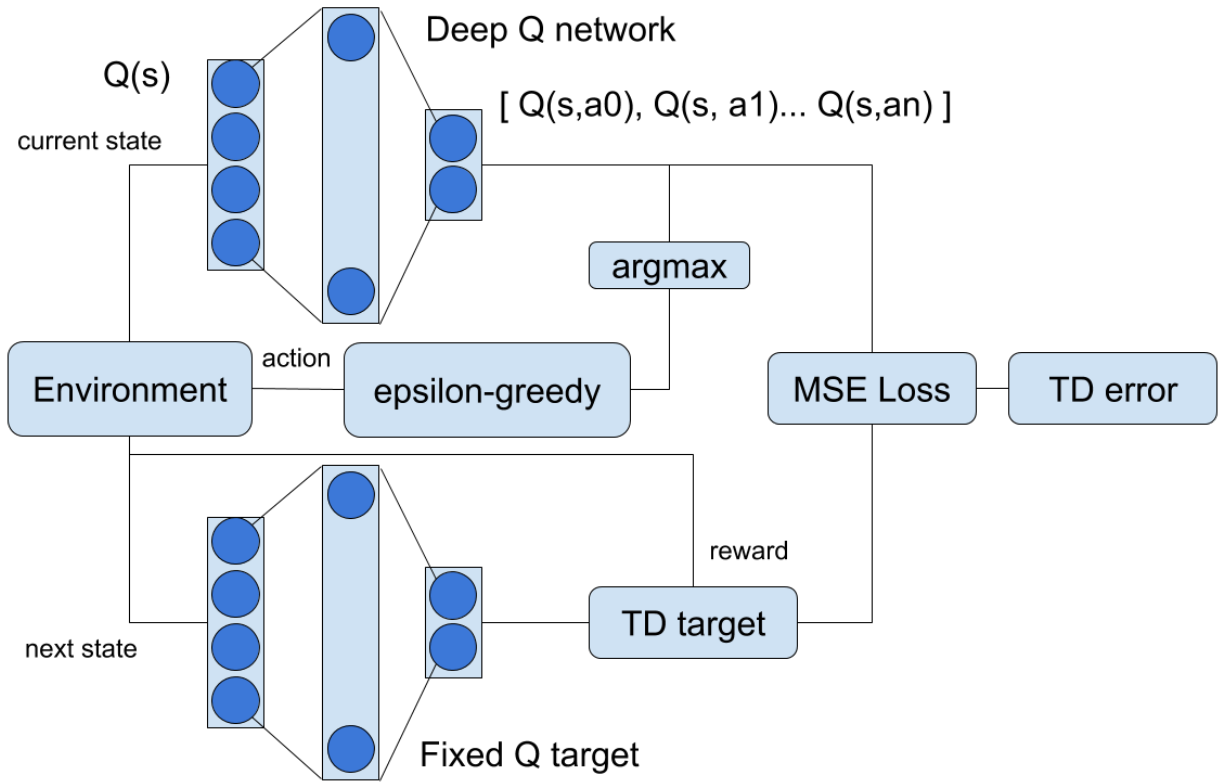


Figure 3: DQN structure and loss

The hidden size of network is 32. To avoid unstable training, use another Q network as target network to fixed target value. And implement in pytorch as follows:

---

```

class DQN(nn.Module):
    def __init__(
        self, observation_space, action_space, device,
        lr=0.05, hidden_size=32, buffer_size=5000
    ):
        super(DQN, self).__init__()

        self.observation_space = observation_space
        if len(observation_space.shape) > 0:
            self.otype = observation_space.shape[0]
        else:
            raise NotImplementedError()

```

```

self.action_space = action_space
if len(action_space.shape) == 0:
    self.asize = action_space.n
else:
    raise NotImplementedError()

self.hidden_size = hidden_size

def __DQN():
    return nn.Sequential(
        nn.Linear(self.otype, self.hidden_size),
        nn.ReLU(inplace=True),
        nn.Linear(self.hidden_size, self.asize, bias=False),
    )

self.device = device
self.eval_net = __DQN().to(self.device)
self.target_net = __DQN().to(self.device)

self.criterion = nn.MSELoss()

```

---

## 4 DDPG structure and loss

DDPG has two network, one is critic and another is actor. Critic is similar as DQN but it get action as input and only output one value. Actor uses current state to find action.

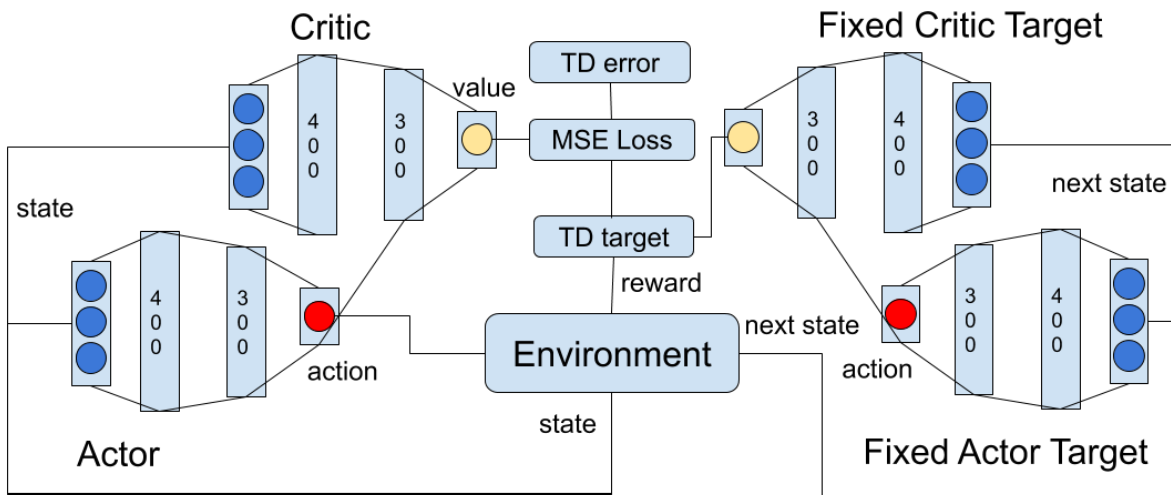


Figure 4: DDPG structure and loss

All layers are fully connect layer. The important point is critic need to feed action into second fully connect layer and actor need use tanh activation to control action range The whole network structure as follows:

---

```
def fanin_init(size, fanin=None):
    fanin = fanin or size[0]
    v = 1. / np.sqrt(fanin)
    return torch.Tensor(size).uniform_(-v, v)

class Actor(nn.Module):
    def __init__(self, osize, asize, init_w=3e-3):
        super(Actor, self).__init__()

        self.osize = osize
        self.asize = asize

        self.main = nn.Sequential(
            nn.Linear(self.osize, 400),
            nn.ReLU(inplace=True),
            nn.Linear(400, 300),
            nn.ReLU(inplace=True),
            nn.Linear(300, self.asize),
            nn.Tanh(),
        )

        self.initweight(init_w)

    def initweight(self, init_w):
        count = 0
        for m in self.main.modules():
            if isinstance(m, nn.Linear):
                count += 1
                m.bias.data.fill_(0.0)
                if count < 3:
                    m.weight.data = fanin_init(m.weight.data.size())
                else:
                    m.weight.data.uniform_(-init_w, init_w)

    def forward(self, x):
        out = self.main(x)
        return out

class Critic(nn.Module):
    def __init__(self, osize, asize, init_w=3e-3):
```

```

    super(Critic, self).__init__()

    self.otype = osize
    self.aseize = asize

    self.main1 = nn.Sequential(
        nn.Linear(self.otype, 400),
        nn.ReLU(inplace=True),
    )

    self.main2 = nn.Sequential(
        nn.Linear(400 + self.aseize, 300),
        nn.ReLU(inplace=True),
        nn.Linear(300, 1),
    )

    self.initweight(init_w)

def initweight(self, init_w):
    count = 0
    for m in list(self.main1.modules()) + list(self.main2.modules()):
        if isinstance(m, nn.Linear):
            count += 1
            m.bias.data.fill_(0.0)
            if count < 3:
                m.weight.data = fanin_init(m.weight.data.size())
            else:
                m.weight.data.uniform_(-init_w, init_w)

def forward(self, x, a):
    out = self.main1(x)
    out = self.main2(torch.cat([out, a.reshape(-1, 1)], 1))
    return out

class DDPG(nn.Module):
    def __init__(
        self, observation_space, action_space
    ):
        super(DDPG, self).__init__()

        self.observation_space = observation_space
        self.action_space = action_space

        if len(observation_space.shape) > 0:
            self.otype = observation_space.shape[0]

```

```

else:
    raise NotImplementedError()

if len(action_space.shape) > 0:
    self.aseize = action_space.shape[0]
else:
    raise NotImplementedError()

# u(s) -> dis of action
self.actor = Actor(self.oseize, self.aseize).to(device)
self.actor_ = Actor(self.oseize, self.aseize).to(device)
self.actor_optim = torch.optim.Adam(self.actor.parameters(),
    ↪ lr=1e-4)

# Q(s, a) -> value
self.critic = Critic(self.oseize, self.aseize).to(device)
self.critic_ = Critic(self.oseize, self.aseize).to(device)
self.critic_optim = torch.optim.Adam(self.critic.parameters(),
    ↪ lr=1e-3)

self.update()

self.memory = Buffer(self.oseize, self.aseize, 10000)
self.store = self.memory.store

self.criterion = nn.MSELoss()

```

---

## 5 How to train DQN

Deep Q Network likes Q-learning that need to estimate  $Q(s, a)$  value in lookup table. And DQN use DNN as lookup table that make DQN can handle continuous state. Therefore DQN update lookup table using back propagation as training model. I used Q target value as label or ground-truth. The equation of target value as follows:

$$Q_{target} = r_t + \gamma \max_a Q(s_{t+1}, a | \theta_Q) \quad (1)$$

If  $s_{t+1}$  is terminated state, then equation is only  $Q_{target} = r_t$ .

### 5.1 Fixed target

But update DQN frequently will make training unstable. Because DNN model not yet convergence and its value is used to next target value. To solve this problem, use fixed target model.

In the begin, create two DQN model, one is target model another is evaluation model.  $Q, \hat{Q}$ . Convert updating equation as follows:

$$Q_{target} = r_t + \gamma \max_a \hat{Q}(s_{t+1}, a | \theta_{\hat{Q}}) \quad (2)$$

Finally update target  $\hat{Q}$  network through evaluation network  $Q$  by every specific steps.

## 5.2 Experience reply

Because state is continuous, its space is very huge. If take whole transition in one game is too correlation and also good transition in one game is rare. Thus I store transition in buffer and take mini batch transition from buffer to train DQN.

---

```
class DQN(nn.Module):
    # Omit ...

    # update target net by evaluation net
    def update(self):
        self.target_net.load_state_dict(self.eval_net.state_dict())

    def store_transition(self, s, a, r, s_next, done):
        transition = np.hstack((s, [a, r], s_next, done))
        # store
        self.buffer[self.buffer_count%self.buffer_size, :] = transition

        self.buffer_count += 1

    def learn(self, batch_size, gamma=0.95):

        # pick data from buffer
        pick_i = np.random.choice(self.buffer_size if self.buffer_count >
        ↪ self.buffer_size else self.buffer_count, size=batch_size)
        x = self.buffer[pick_i, :]

        # clear grad
        self.eval_optimizer.zero_grad()

        # only use state
        q_eval, q_target = self( torch.tensor( x[:, :self.otype]
        ↪ ).to(self.device) , torch.tensor( x[:, -(self.otype+1):-1]
        ↪ ).to(self.device))

        # x[:, -1] done list
        q_target_value = torch.max(q_target, dim=1)[0].detach()
        q_target_value[np.argwhere(x[:, -1] == 1).reshape(-1)] = 0.0
```



```

y = q_eval.clone().detach()
# set y = r + \gamma max_a( Q_target(s', a) )
y[np.arange(batch_size), x[:, self.oseize]] = torch.tensor(x[:,
    ↪ self.oseize+1]).float() + (gamma * q_target_value)

# get loss
loss = self.criterion(q_eval, y)

# update weight
loss.backward()

# prevent too huge grad
#for param in self.eval_net.parameters():
#    param.grad.data.clamp_(-1, 1)

self.eval_optimizer.step()

return loss

```

---

## 6 How to implement epsilon-greedy action

If always use action result from model, model can't explore new space in environment. To avoid this problem, model sometimes random a action.

---

```

epsilon_thr = epsilon_end + ( (epsilon_start - epsilon_end) * np.exp(-1. *
    ↪ epoch / epsilon_decay) )

# choose action
if np.random.rand() < epsilon_thr:
    action = np.random.randint(agent.aseize)
else:
    action = agent.think(observation)

```

---

I decrease epsilon threshold along with epoch increasing.

## 7 Critic and Actor updating in DDPG

DDPG also use fixed target and experience reply to train model. I initial actor as  $u(s|\theta_u)$  and  $\hat{u}(s|\hat{\theta}_u)$  and critic as  $Q(s, a|\theta_Q)$  and  $\hat{Q}(s, a|\hat{\theta}_Q)$ . Critic likes DQN, it need to estimate value. Thus update it with equation as follows:

$$Q_{target} = r_t + \gamma \hat{Q}(s_{t+1}, \hat{u}(s_{t+1}|\theta_{\hat{u}})|\theta_{\hat{Q}}) \quad (3)$$

Loss can use mean square error as follows:

$$L = \frac{1}{N} \sum (Q_{target} - Q(s_t, a_t|\theta_Q))^2 \quad (4)$$

Using previous equation let critic can know how many reward in specific state and action. Because of it, actor depend on critic to know which action to generate better reward.

## 7.1 How to calculate the actor's gradients

Actor's objective is making better reward in continuous action space. I use equation as follows to reach it.

$$L = -Q(s, u(s|\theta_u)|\theta_Q) \quad (5)$$

I minimize loss like maximize reward from actor's action.

Use chain rule to show how to calculate gradient.

$$L = -Q(s, a|\theta_Q), \quad a = u(s|\theta_u) \quad (6)$$

$$\frac{\nabla L}{\nabla \theta_u} = - \frac{\nabla Q(s, a|\theta_Q)}{\nabla a} \frac{\nabla a}{\nabla u(s|\theta_u)} \frac{\nabla u(s|\theta_u)}{\nabla \theta_u} \quad (7)$$

$$= - \frac{\nabla Q(s, a|\theta_Q)}{\nabla u(s|\theta_u)} \frac{\nabla u(s|\theta_u)}{\nabla \theta_u} \quad (8)$$

---

```

class DDPG(nn.Module):
    # Omit ...
    def update(self, tau=None):
        if tau is None:
            self.actor_.load_state_dict(self.actor.state_dict())
            self.critic_.load_state_dict(self.critic.state_dict())
        else:
            soft_update(self.actor_, self.actor, tau)
            soft_update(self.critic_, self.critic, tau)

    def learn(self, batch_size, gamma=0.99):
        s, a, r, s_, done = self.memory.miniBatch(batch_size)

        with torch.no_grad():
            q_ = self.critic_(
                to_tensor(s_),
                self.actor_(to_tensor(s_))
            )

```

```

q_target = to_tensor(r) + ( gamma * to_tensor(1.0 -
    ↪ done.astype(np.float))*q_ )
q_target = q_target.detach()

# Critic learn
self.critic_optim.zero_grad()

#s = to_tensor(s)
q = self.critic(
    to_tensor(s),
    to_tensor(a)
)
loss_value = self.criterion(q, q_target)
loss_value.backward()

self.critic_optim.step()

# Actor learn
self.actor_optim.zero_grad()

loss_policy = -self.critic(
    to_tensor(s),
    self.actor(to_tensor(s))
)

loss_policy = loss_policy.mean()
loss_policy.backward()

self.actor_optim.step()

self.update(tau=1e-3)

return loss_value, loss_policy

```

---

Finally, I need to update target with  $\tau$ .

$$\theta_{\hat{Q}} = (1 - \tau)\theta_{\hat{Q}} + \tau\theta_Q\theta_{\hat{u}} = (1 - \tau)\theta_{\hat{u}} + \tau\theta_u \quad (9)$$

---

```

def soft_update(target, source, tau):
    for target_param, param in zip(target.parameters(),
    ↪ source.parameters()):
        target_param.data.copy_(

```

```
        target_param.data * (1.0 - tau) + param.data * tau
    )
```

---

## 8 Code explanation

First, create environment. I wrap open AI gym can show in jupyter notebook.

---

```
def myrender(self, title=''):
    plt.title(title)
    plt.imshow(self.old_render(mode='rgb_array'))
    plt.axis('off')
    plt.show()
    clear_output(wait=True)
    time.sleep(1/30)

def wrapper_gymenv(env):
    env = NormalizedAction(env) # only need in DDPG
    env.old_render = env.render
    env.render = lambda title='': myrender(env, title)
    return env

env = gym.make('Pendulum-v0')
# or env = gym.make('CartPole-v0')
env = wrapper_gymenv(env)
```

---

Second, create agent.

---

```
agent = DDPG(env.observation_space, env.action_space)
# or agent = DQN(env.observation_space, env.action_space,
→ torch.device("cpu"))
```

---

Third, training agent with environment. In DQN, I need to store best agent in training.

---

```
train(agent, 3500, 64)
# best_agent = train(agent, 1000, 128, 50) in DQN
```

---

Train procedures are very similar between DQN and DDPG. I use DDPG training as example. In every epoch, I reset environment and send action to environment through agent or random action. (In DQN, use epsilon-greedy action) When environment responses next state and reward, I store state, action, reward, next state and done (Is game end ?) into experience buffer. If buffer has enough transition, I let agent learn from experience buffer

with mini batch size. (In DDPG case, it is 64) If environment reach terminated state, I evaluation agent to get average total rewards. After that I store some metric and show result.

---

```
def train(
    agent, epoch_size, batch_size
):

    for epoch in range(len(agent.metrics), epoch_size):
        # play

        observation = env.reset()
        score = 0.0
        step = 0

        loss_v_t = []
        loss_p_t = []

        agent.train()
        agent.reset()

        while True:

            if agent.memory.buffer_count <= 100:
                action = agent.random_think()
            else:
                action = agent.think(observation, True)

            # do action
            observation_next, reward, done, info = env.step(action)
            step += 1

            agent.store(observation, action, reward, observation_next, done)

            if agent.memory.buffer_count > 100:
                loss_value, loss_policy = agent.learn(batch_size)
                loss_v_t.append(loss_value.item())
                loss_p_t.append(loss_policy.item())

            # update state
            observation = observation_next

            if done:
                break
```

---

```

test_score = np.mean([play(agent, show=False) for _ in range(100)])

# metrics
loss_v_t = np.average(loss_v_t) if len(loss_v_t) > 0 else 0.0
loss_p_t = np.average(loss_p_t) if len(loss_p_t) > 0 else 0.0

agent.metrics.append( (test_score, loss_v_t, loss_p_t) )

clear_output(wait=True)
print('#{:6.0f}: score : {:5.4f}'.format(epoch, test_score))

```

---

Fourth, evaluation agent without any random.

---

```

def evaluation(
    agent, eval_size
):
    scores = []
    for i in range(eval_size):
        scores.append( play(agent, show=False) )

    print('mean = {}, max = {}'.format(np.mean(scores), np.max(scores)))

    return scores

scores = evaluation(agent, 1000)

```

---

## 9 CartPole with DQN performance

mean	max	min	median	count	max step
200	200	200	200	100	200

Table 1: CartPole with DQN performance

## 10 Pendulum with DDPG performance

mean	max	min	median	count	max step
-153.861	-2.256	-533.589	-123.861	1000	100

Table 2: Pendulum with DDPG performance

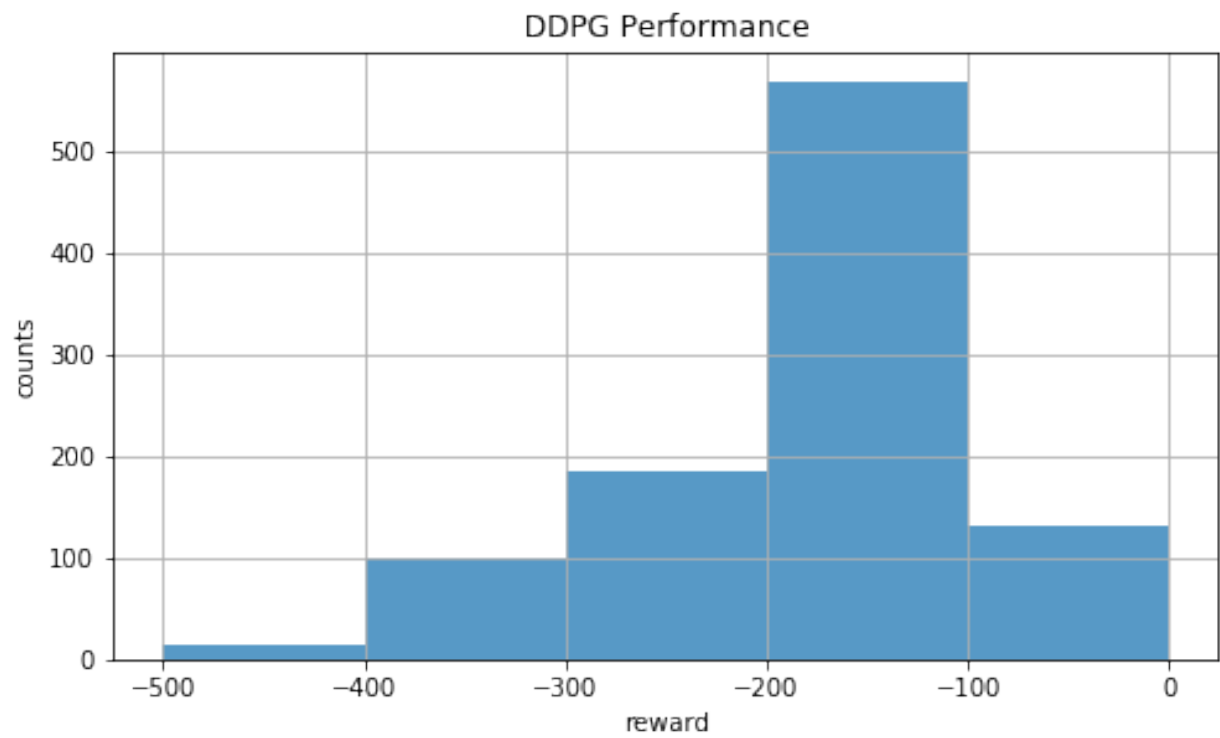


Figure 5: DDPG rewards distribution